

DSM

Note Title

07-Dec-22

- False sharing

- FT: Checkpoints: Inconsistency

for
 $x++$
 $y \leftarrow y+x$


for

$x \leftarrow x+1$

for


$y \leftarrow y+1$

$\boxed{y^x}$



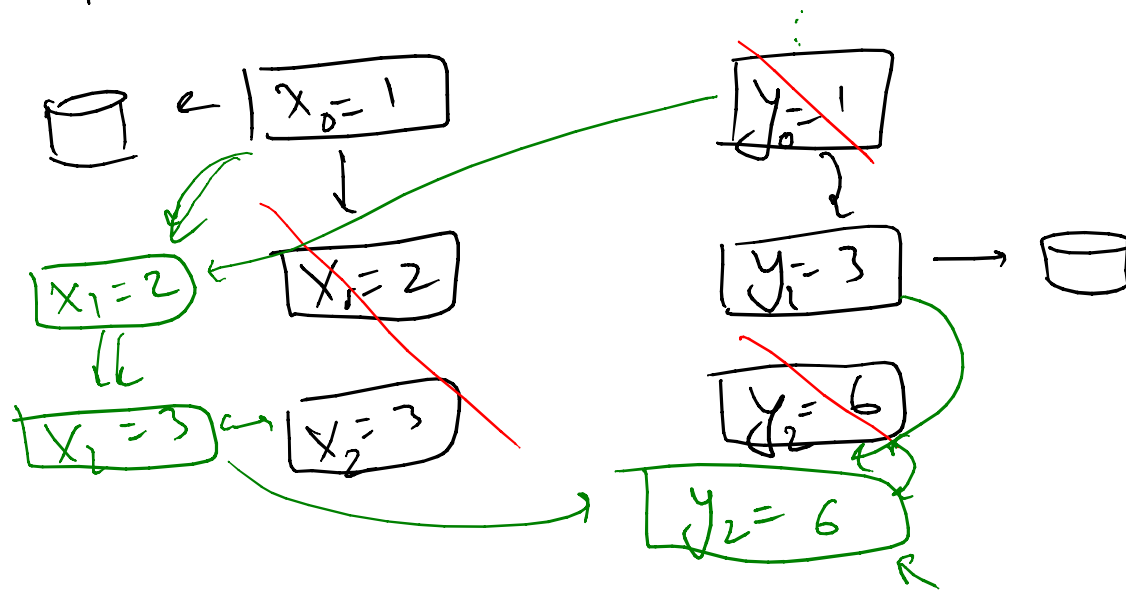
$\boxed{x=2}$ $\boxed{y=5}$

Replication - Slow



MapReduce Locality^L / FT → deterministic ✓
Idempotence

Inconsistent checkpoints

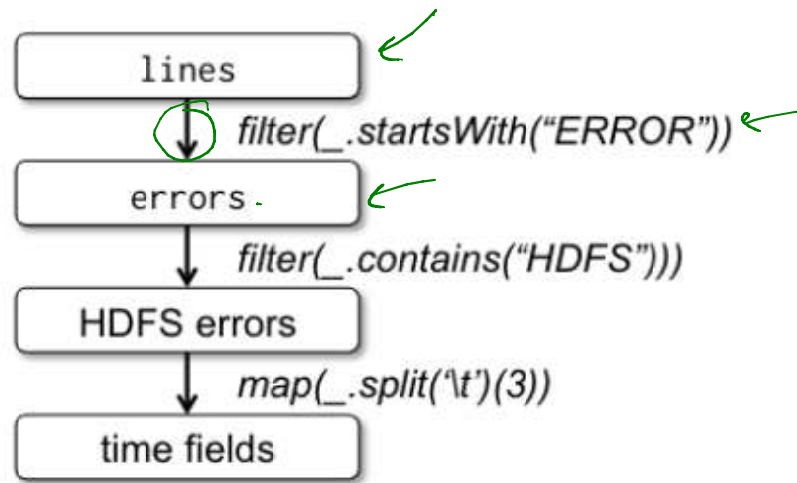


→ Immutable ←

for
 → $x++$
 $y \leftarrow y+x$

$X_0=1 \leftrightarrow X_0=1$

→ Lineage ←





Resilient
Distributed
Datasets

Figure 1: Lineage graph for the third query in our example. Boxes represent RDDs and arrows represent transformations.

```

{
  lines = spark.textFile("hdfs://...")
  errors = lines.filter(_.startsWith("ERROR"))
  errors.persist()
}
  
```






RDDs — partitions

P — worker

Page Rank




0.15
0.05

u_1

$$u_1, u_3 < u_2$$

$$u_1, u_3 < u_4$$

$$u_2 < u_4$$



lines

u1, u2
u3, u2
u2, u4
u1, u1
u2, u2
u3, u3
u4, u4

links

u1	{ u1, u2 }
u2	{ u2, u4 }
u3	{ u3, u2 }
u4	{ u4 }

HDFS

rank_{u1} 0

u1	1
u2	1
u3	1
u4	1

contrib



u1	0.5
u2	0.5
u2	0.5
u4	0.5
u3	0.5
u2	0.5
u4	1

rank₁

u1	$0.5 * 0.85 + 0.15 = 0.575$
u2	$1.5 * 0.85 + 0.15 = 1.425$
u3	$0.5 * 0.85 + 0.15 = 0.575$
u4	$1.5 * 0.85 + 0.15 = 1.425$

HDFS

HDFS



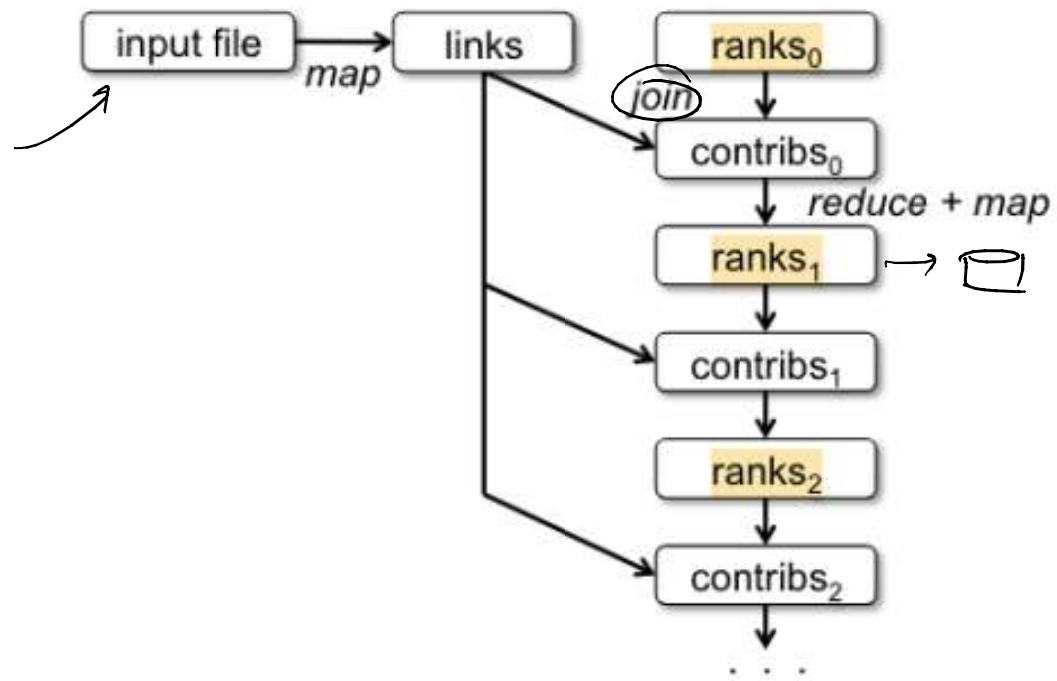


Figure 3: Lineage graph for datasets in PageRank.

125M

Am fine grained

Consistency

Checkpoint

Spark

Deterministic coarse

Immutable

FT: Reexec ← →

MR