

iReplayer: In-situ and Identical Record-and-Replay for Multithreaded Applications

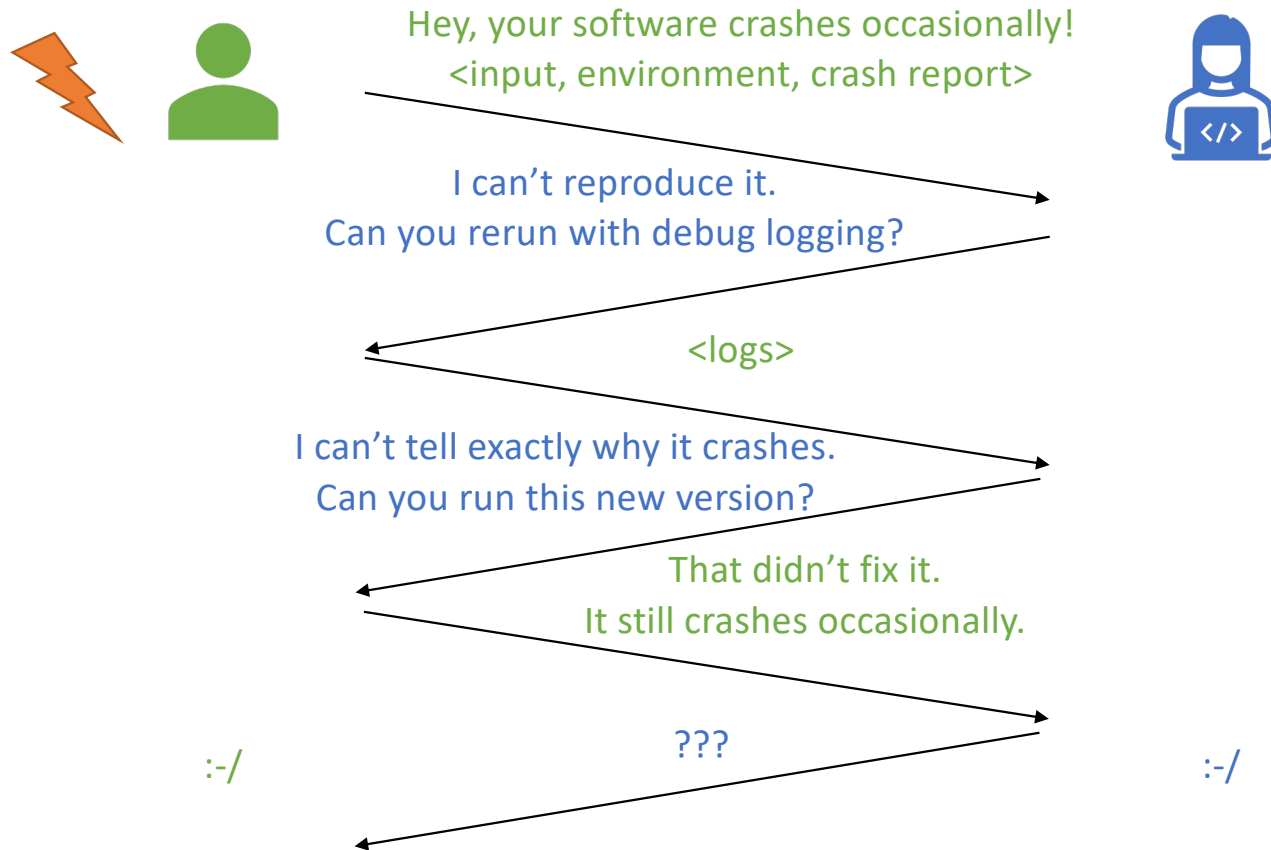
Hongyu Liu, Sam Silvestro, Wei Wang, Chen Tian*, Tongping Liu

University of Texas at San Antonio, Huawei US Lab*

PLDI 2018

Discussion Lead: Abhilash Jindal

Why record-and-replay?



Why is reproducing hard?

```
def transfer(src: Account, dst: Account, amount: int):  
    if src.bal > amount:  
        src.bal -= amount  
        dst.bal += amount
```

Initial state: src.bal = dst.bal = 800

T1:
Amount: 500
R src.bal -> 800
W src.bal -> 300
W dst.bal -> 1300

T2:
Amount: 400
R src.bal -> 300

T1:
Amount: 500
R src.bal -> 400

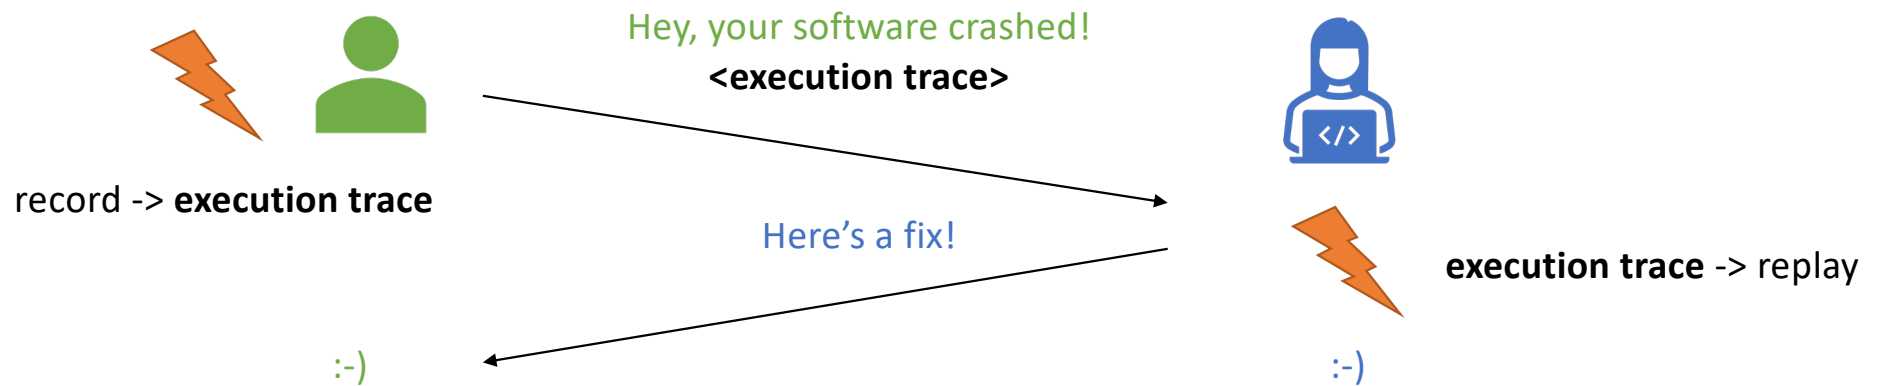
T2:
Amount: 400
R src.bal -> 800
W src.bal -> 400
W dst.bal -> 1200

T1:
Amount: 500
R src.bal -> 800
W src.bal -> 300
W dst.bal -> 1300

T2:
Amount: 400
R src.bal -> 800
W src.bal -> 400
W dst.bal -> 1200

Different executions produce different states

Identical record-and-replay



- Address all sources of non-determinism
- Always on: Low overhead
- Don't leak privacy sensitive information

Why in-situ?




record -> **execution trace**



replay **execution trace** in *same machine process* to debug

- Users don't want to share privacy sensitive trace
 - *But, in-situ doesn't help. Users are not technical enough to debug. Programmers must still see the trace for recovery.*
- Can do automatic online recovery
 - *No evidence shown in the paper to perform such recovery automatically*

Why in-situ? (2)



replay **execution trace** in *same machine process* to debug

- Programmer *is* the user
 - Large scale concurrency testing
 - Production environment monitoring

Easier to do *identical* replay in-situ with low-overhead

- Same hardware
 - Identical binary
 - Identical floating-point behavior
- Same OS
 - Identical memory layout
 - Identical file system
- Same process
 - Identical thread ids: Memory allocators use thread ids





Figure 1. A null reference problem.

Existing RnR Systems

- Offline RnR: replay occurs after recording
 - Instrumentation: iDNA[VEE'06], PinPlay[CGO'10] **10-100X**
 - Offline Assisted Analysis: ODR[sosp'19], CLAP[PLDI'13], Light[PLDI'15], H3[ATC'17], Castor[ASPLOS'17]
Low overhead, but substantial time of offline analysis
 - Custom Hardware: Strata[ASPLOS'06], DeLorean[ISCA'08], Capo[ASPLOS'09] **Impractical**
 - Hybrid Analysis: Chimera[PLDI'12] **40% for 4 threads and hide failures**
- Online RnR: record and replay execute concurrently
 - Speculation Based: Respec[ASPLOS'10] **55% for 4 threads**
 - Uniparallelism: DoublePlay[ASPLOS'11] **28% for 4 threads**
 - iReplayer (this paper)
Unmodified OS, hardware, compiler. Low overhead

iReplayer: Recording Overhead



iReplayer demo: Interactive Debugging



iReplayer demo: Interactive Debugging



iReplayer demo: Interactive Debugging



iReplayer demo: Interactive Debugging



iReplayer demo: Interactive Debugging



iReplayer demo: Interactive Debugging



iReplayer demo: Interactive Debugging



iReplayer demo: Interactive Debugging




iReplayer demo: Interactive Debugging



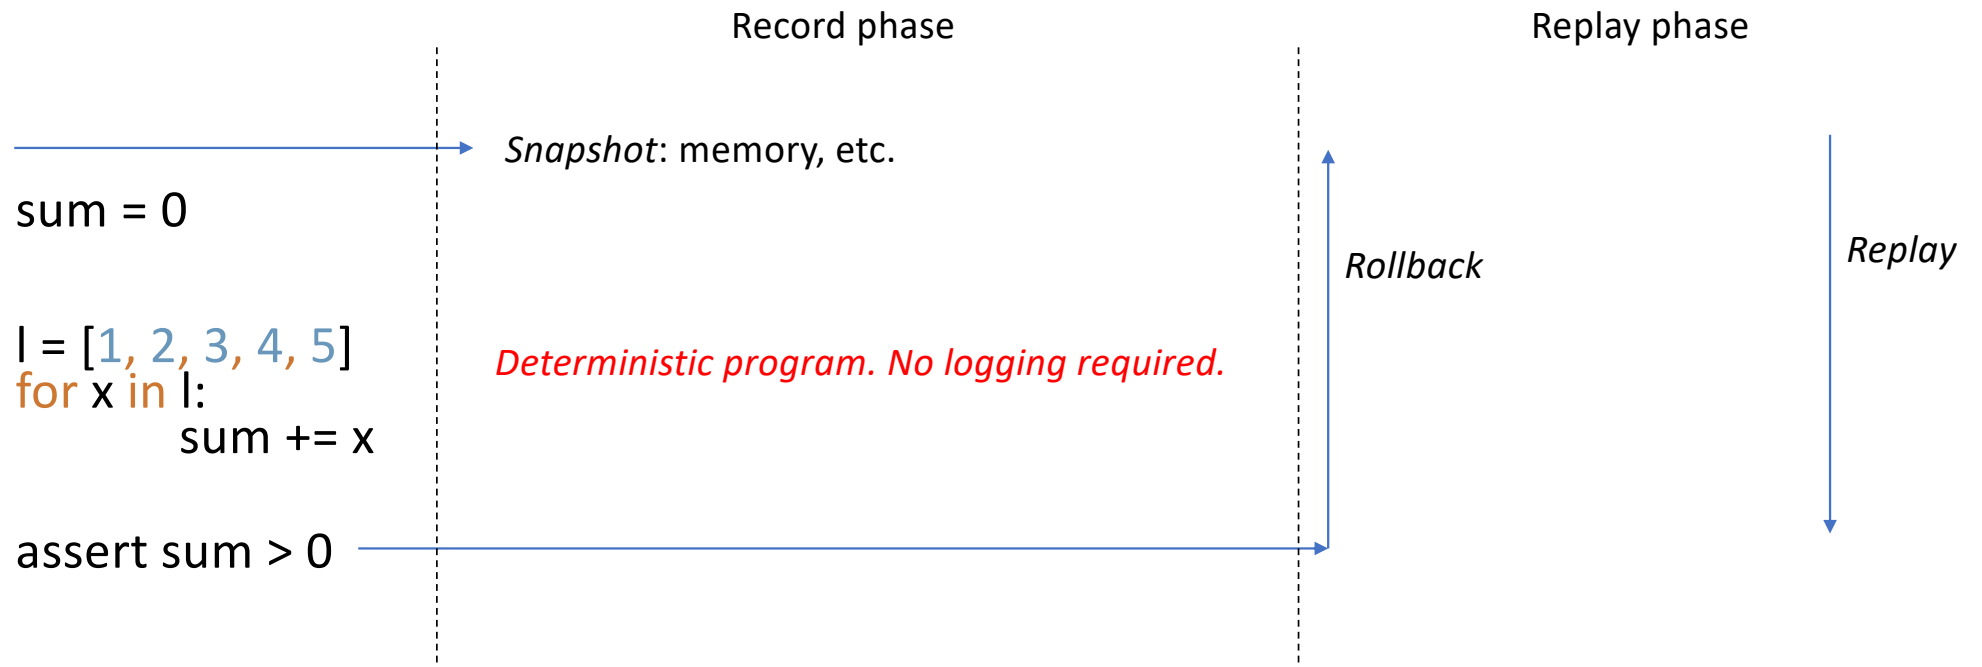
iReplayer demo: Interactive Debugging



Basic Idea



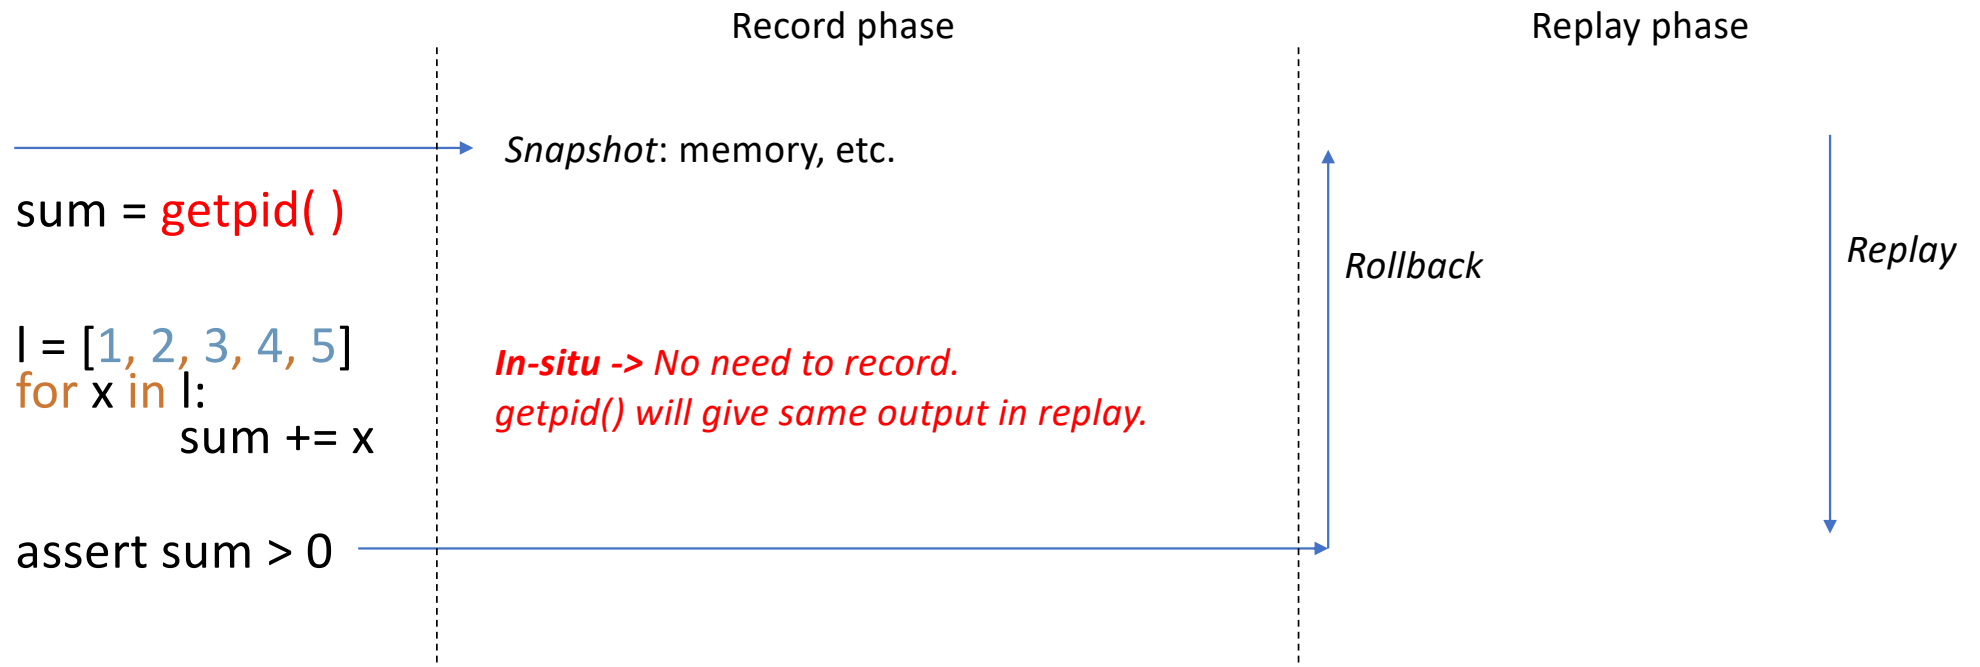
Basic Idea



Addressing sources of non-determinism

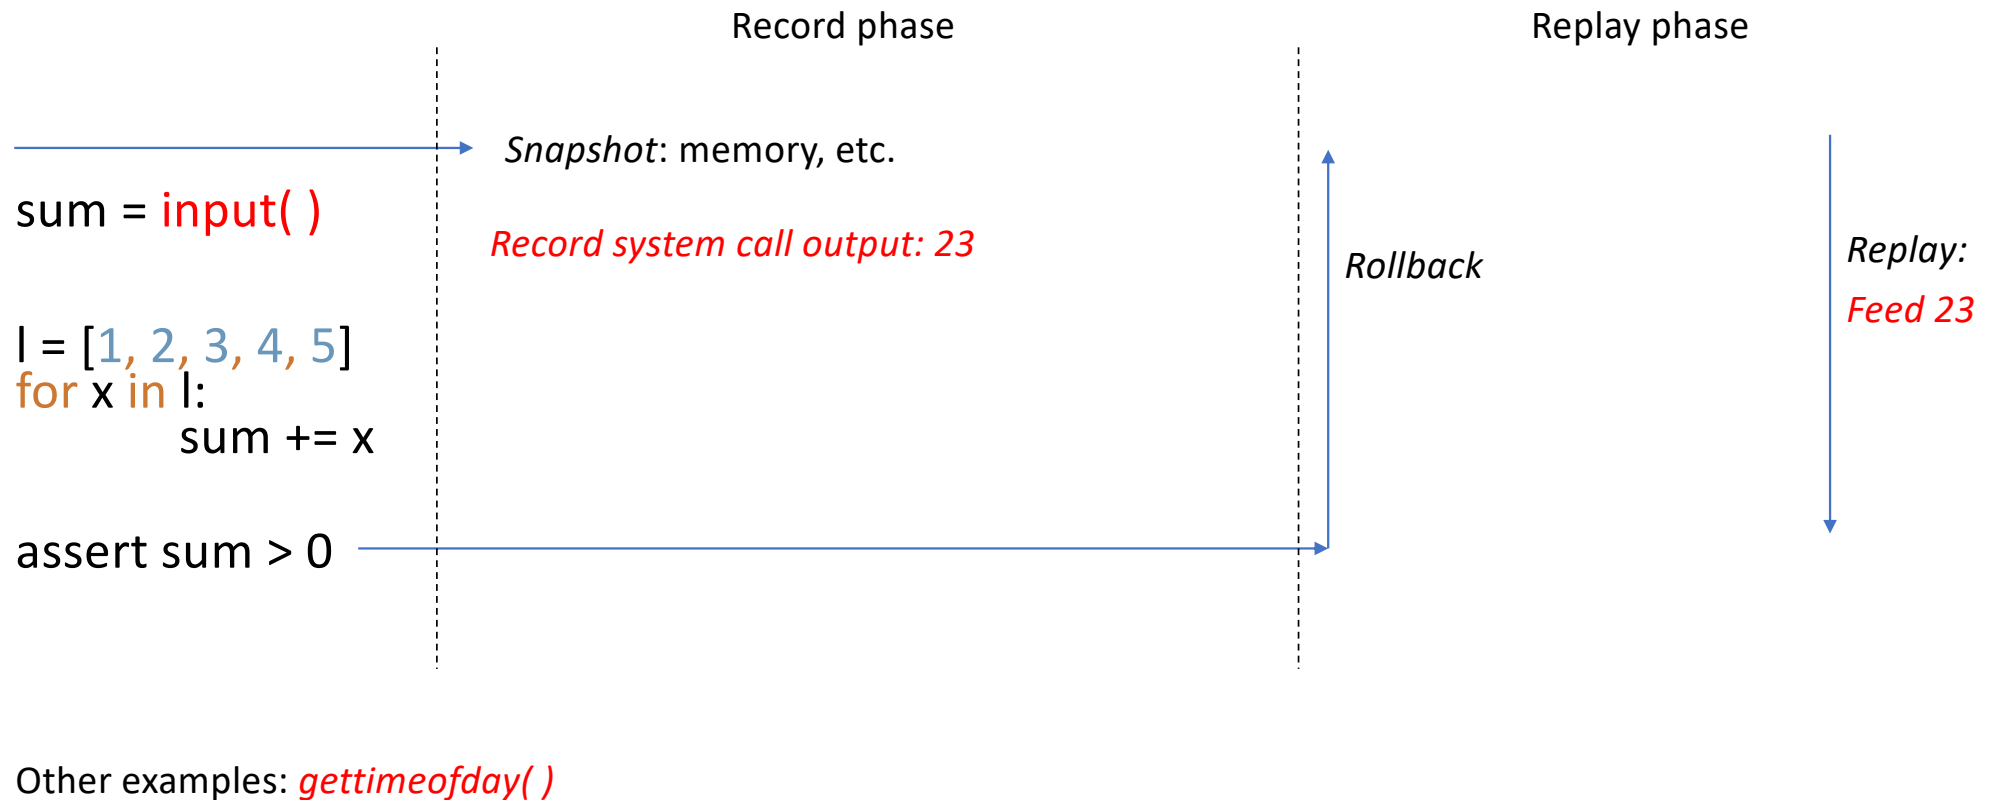
- **Single thread**
 - **System calls: Repeatable, recordable, revocable, deferrable, irrevocable**
- Multithreading
 - Thread lifecycle
 - Memory allocators
 - Synchronization
 - Racy memory accesses

Repeatable system calls

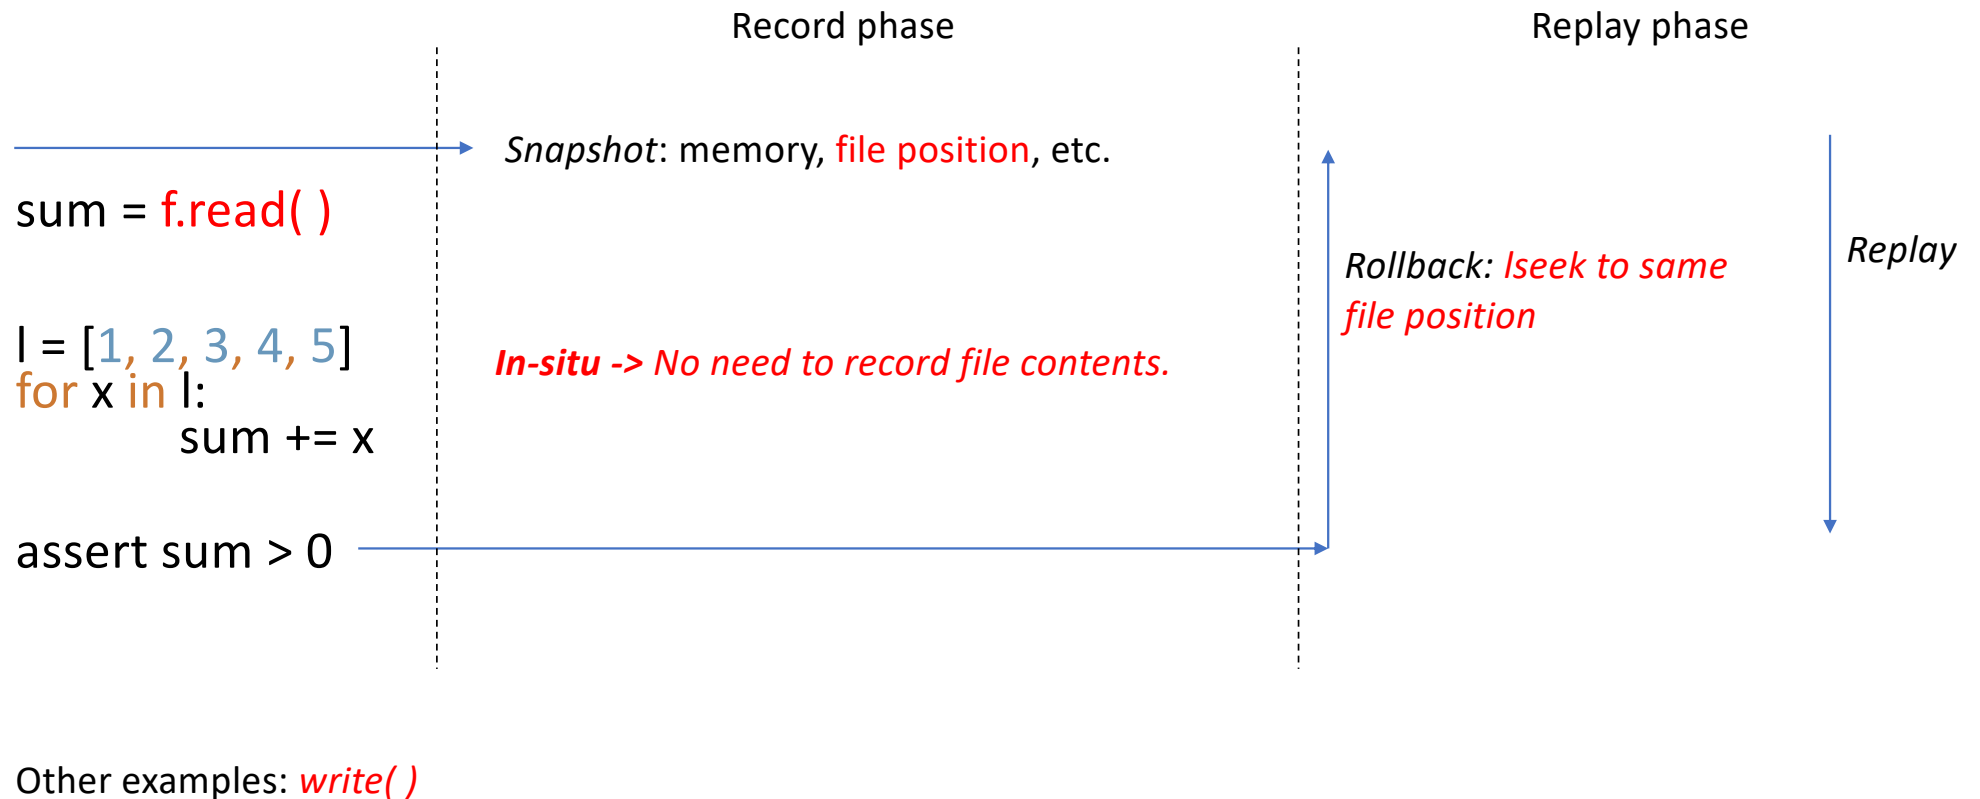


Other examples: *getcwd()*

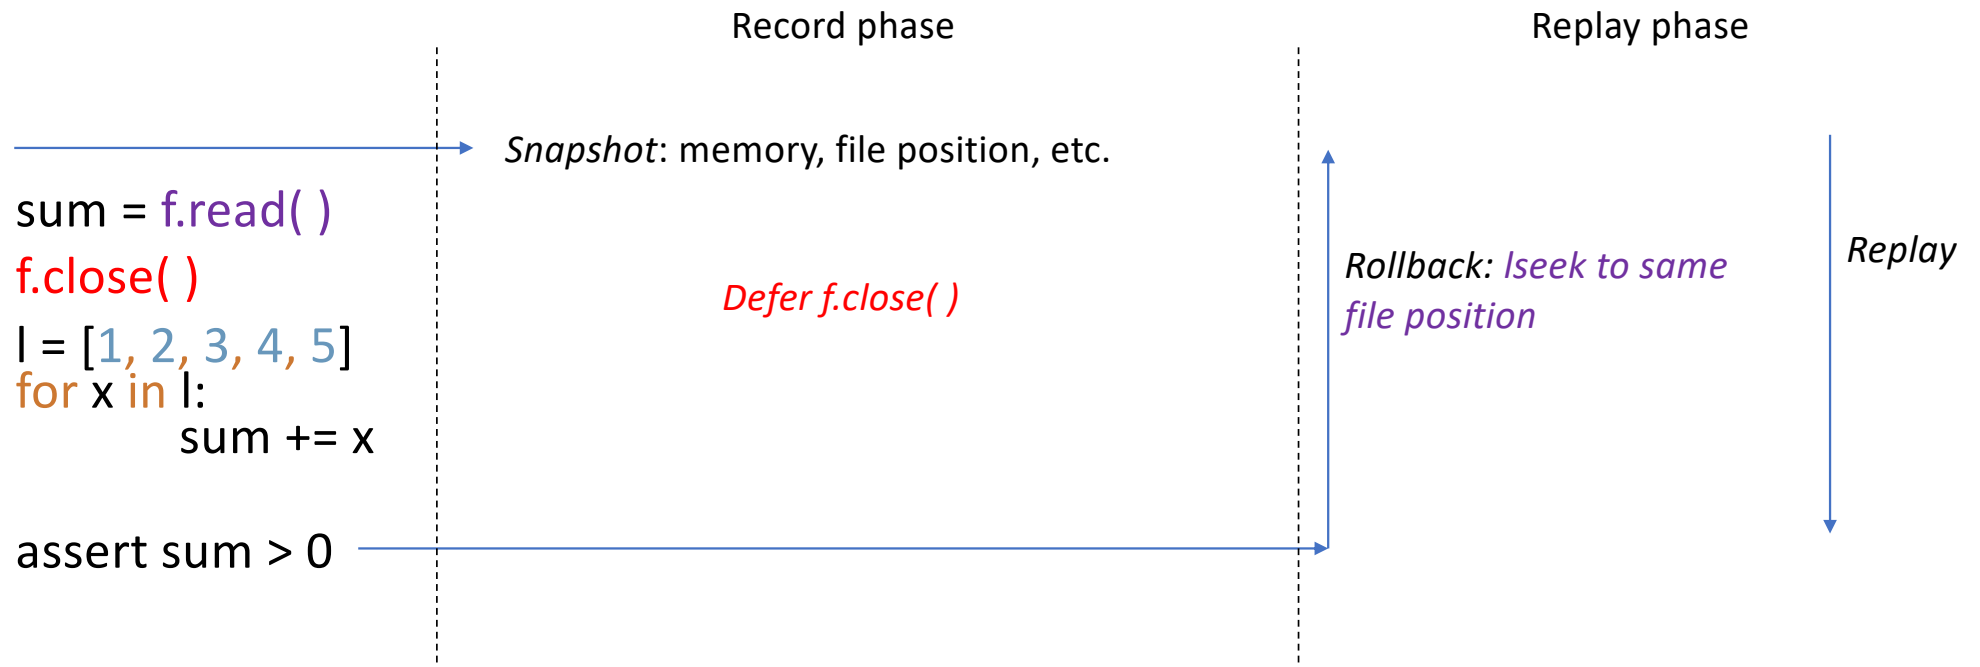
Recordable system calls



Revocable system calls

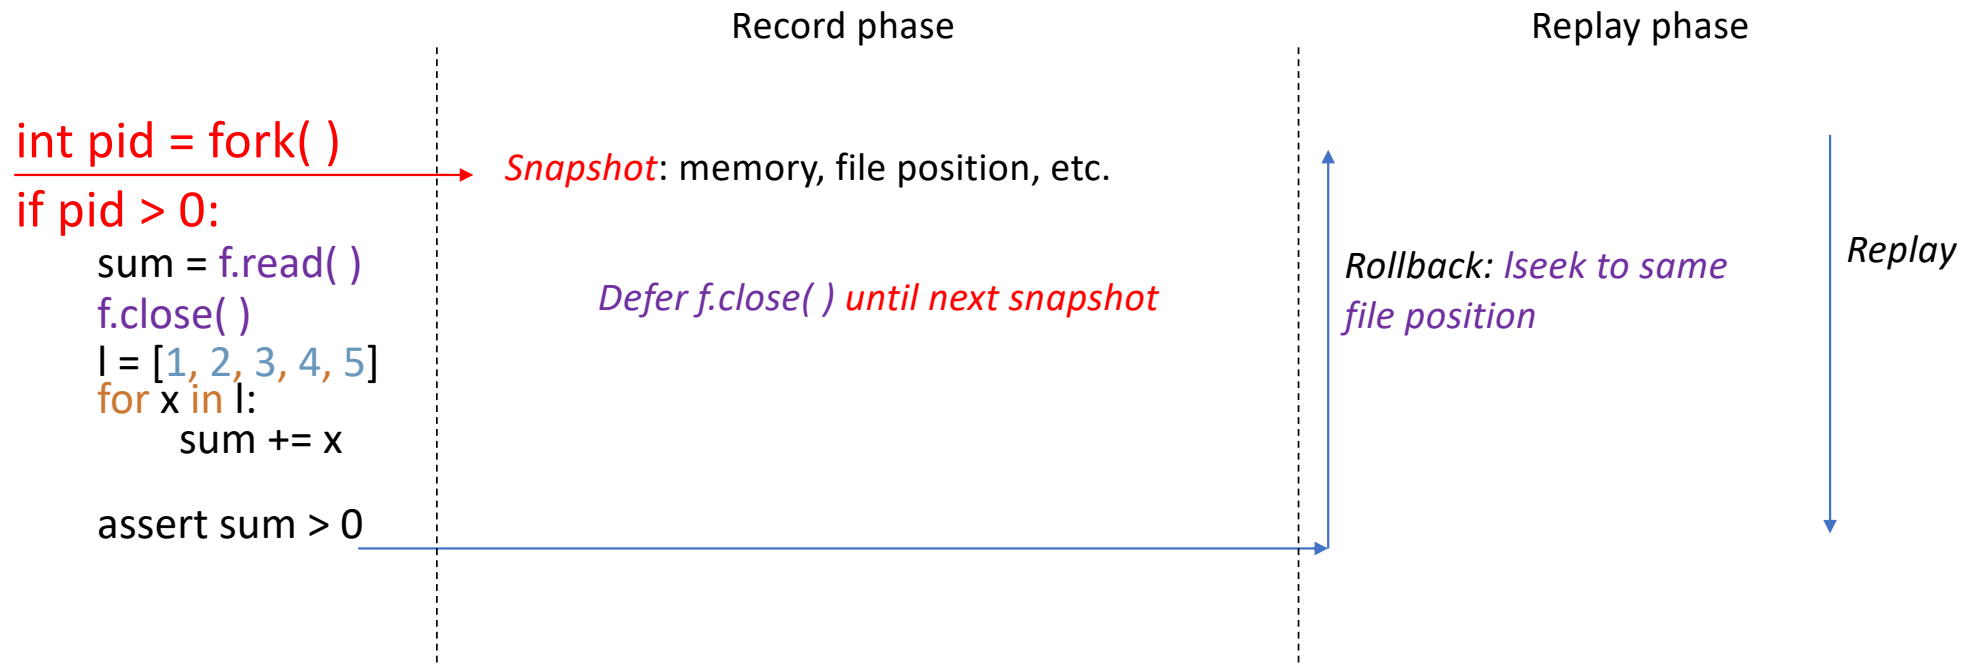


Deferrable system calls




Other examples: *socket.close()*, *munmap*

Irrevocable system calls



Don't know how to rollback irrevocable syscalls. So, we won't allow rolling back to a state prior to that syscall.

Epoch-based record replay



Cannot rollback to a state prior to the current epoch!
Justification: Root cause of bugs is typically not too far from the actual bug

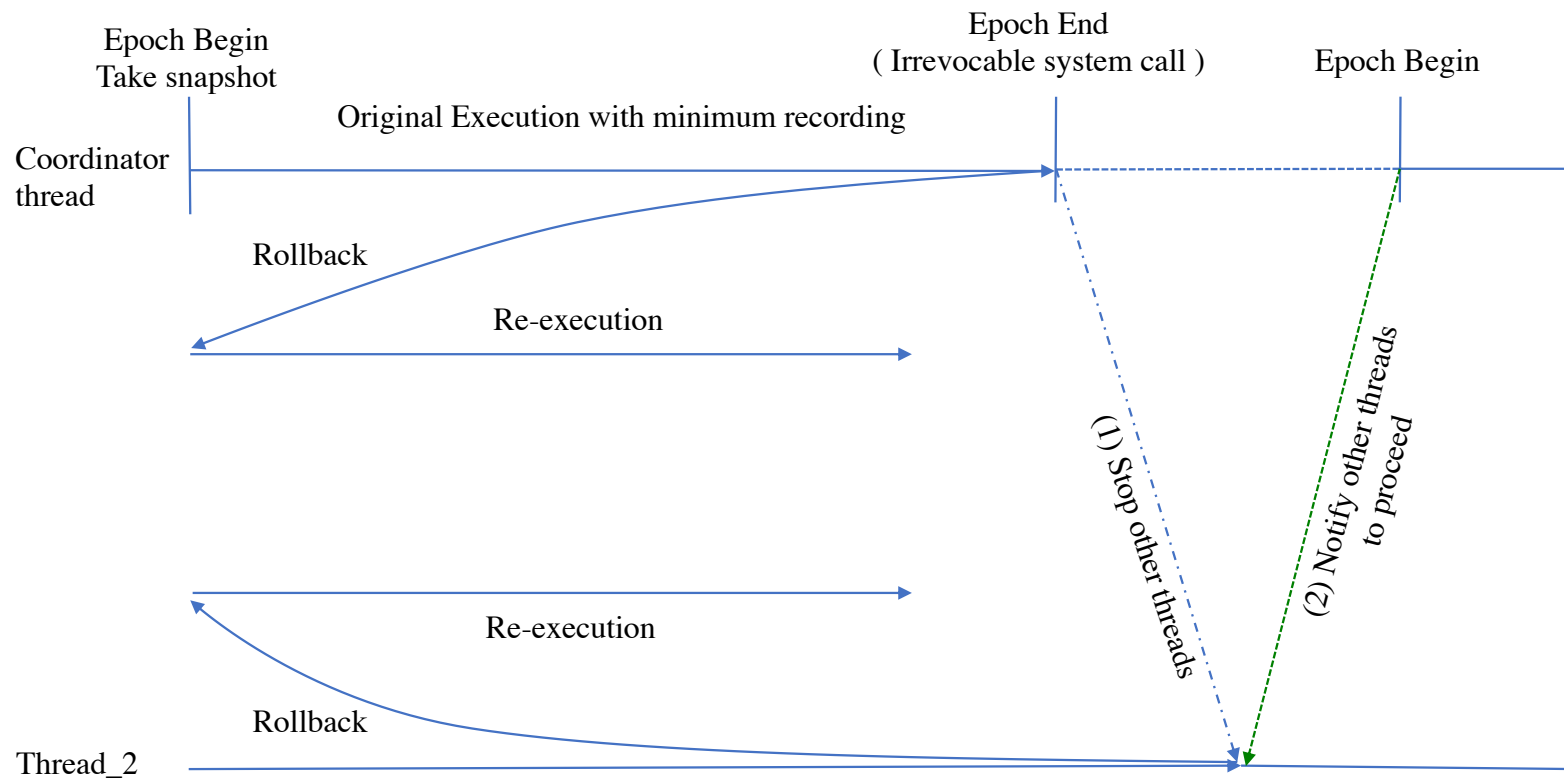
Syscalls: Adapt to In-Situ Setting

Category	Syscall Examples	Handling of the syscall
Repeatable	getpid, getcwd	no handling
Recordable	gettimeofday, mmap, open	record/replay
Revocable	file read/write	rollback side-effect with low overhead
Deferrable	close, munmap, (thread exits)	defer to next epoch
Irrevocable	fork, lseek	stop current epoch

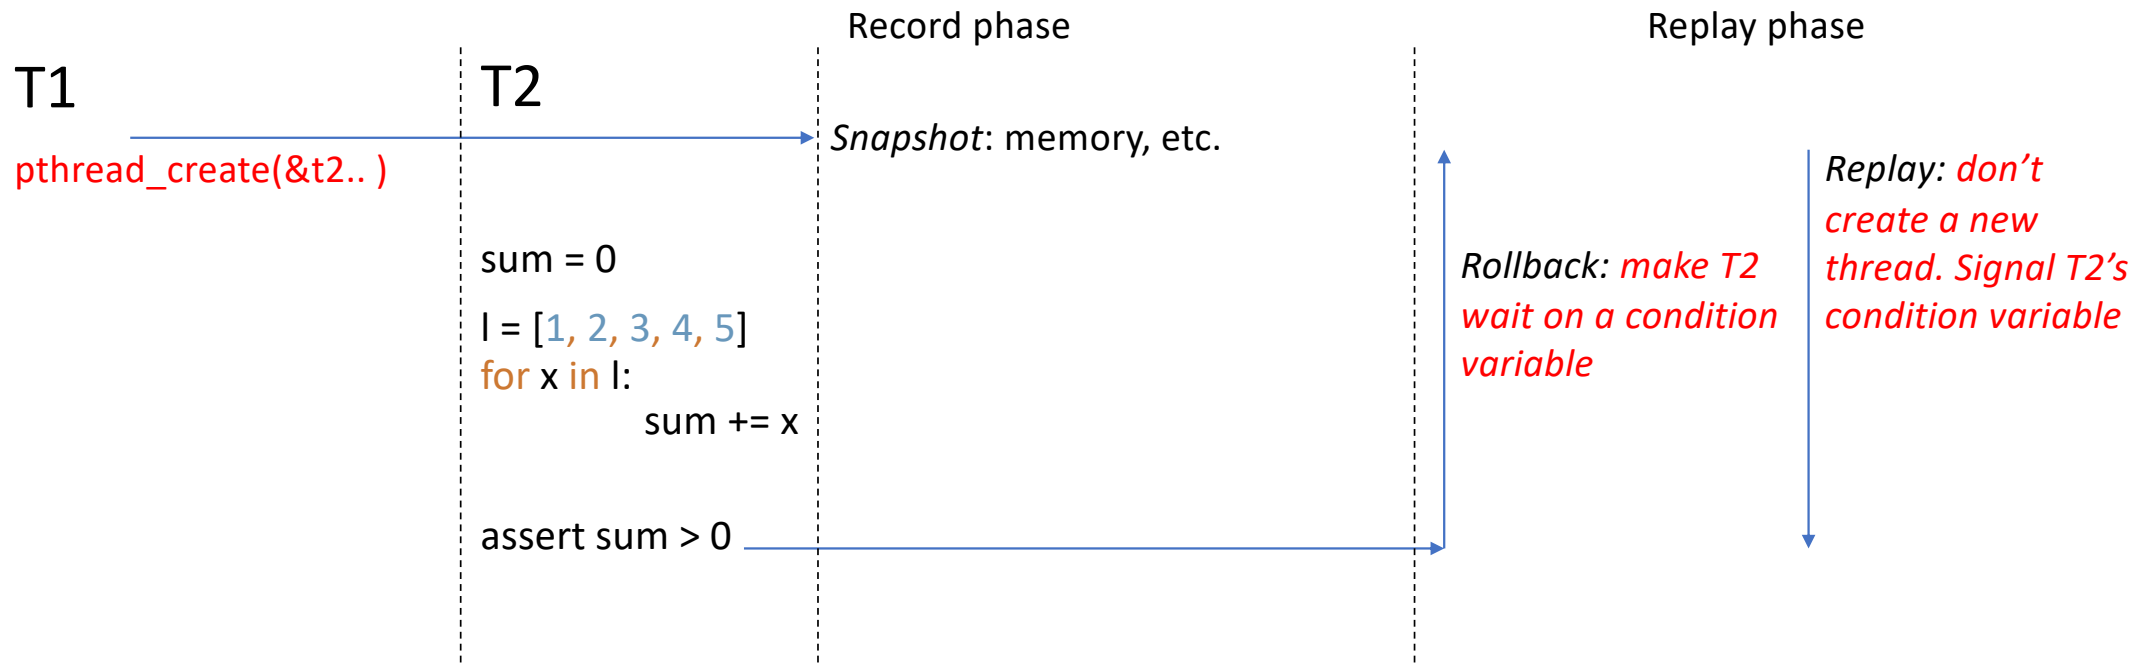
Addressing sources of non-determinism

- Single thread: system calls
 - Repeatable, recordable, revocable, deferrable, irrevocable
- **Multithreading**
 - **Thread lifecycle**
 - **Synchronization**
 - Memory allocators
 - Racy memory accesses

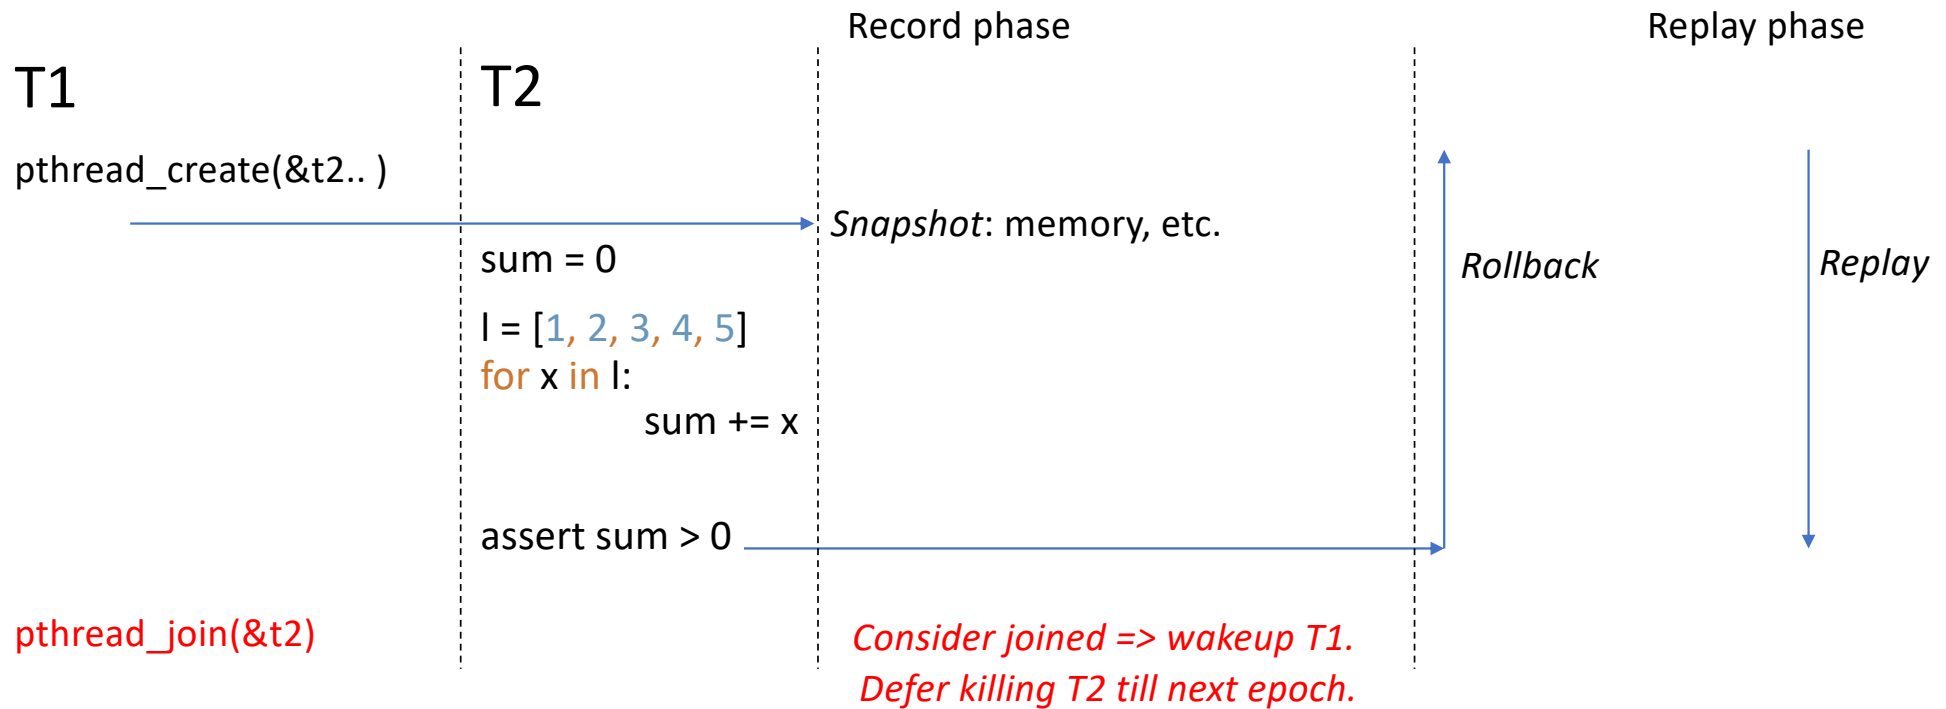
Multithreading



Thread lifecycle

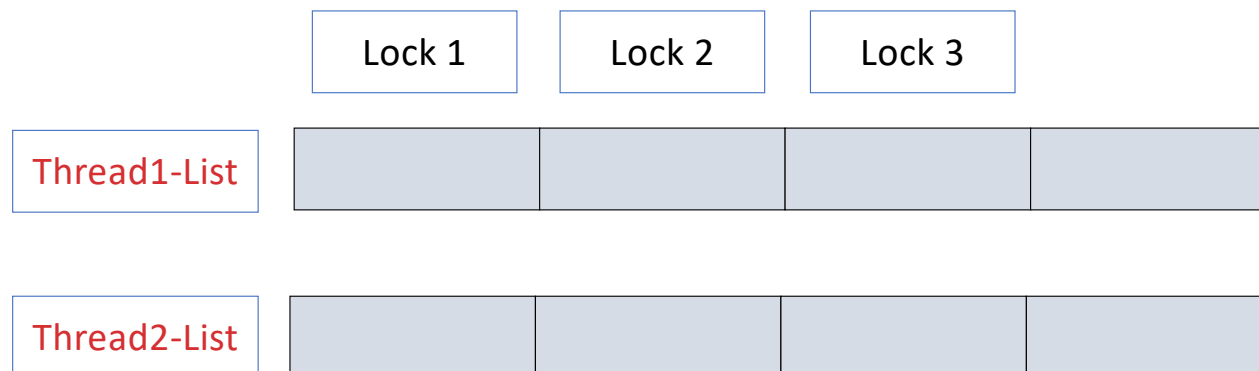


Thread lifecycle (2)



Recording Synchronizations

```
Thread1:      Thread2 :
              Lock (&lock2);
Lock (&lock1);  Unlock (&lock2);
Lock (&lock2);  ←
Unlock (&lock2);
Unlock (&lock1); ←
Lock (&lock3);
Unlock (&lock3); Lock (&lock1);
                  Unlock (&lock1);
```




Recording Synchronizations

```
Thread1:                                Thread2 :
                                           Lock(&lock2);
Lock(&lock1);                             Unlock(&lock2);
Lock(&lock2);                             Lock(&lock1);
Unlock(&lock2);                           Unlock(&lock1);
Unlock(&lock1);
Lock(&lock3);
Unlock(&lock3);
```



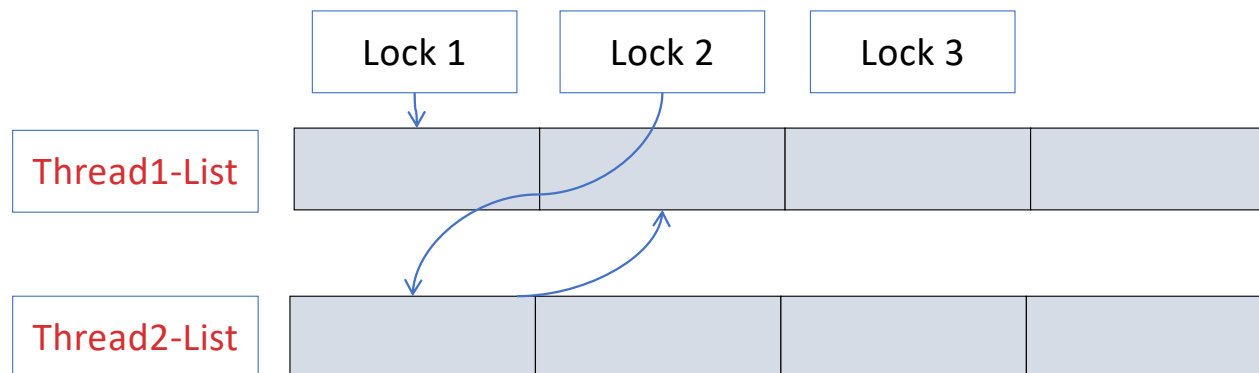
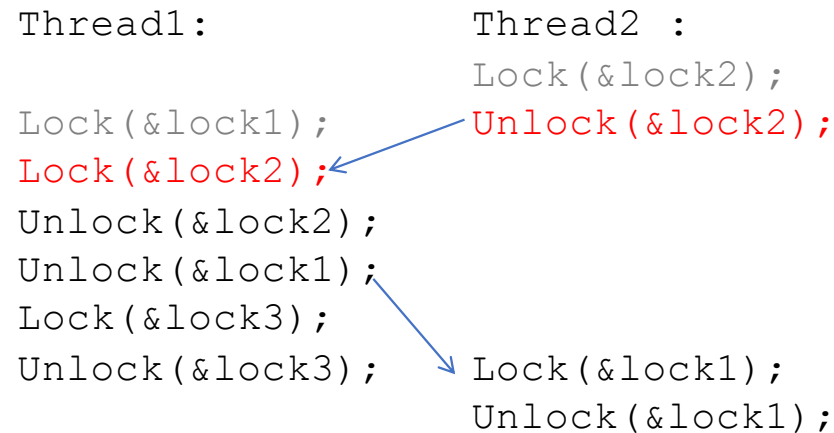
Recording Synchronizations

```
Thread1:      Thread2 :  
              Lock(&lock2);  
Lock(&lock1);  Unlock(&lock2);  
Lock(&lock2);    
Unlock(&lock2);  
Unlock(&lock1);  
Lock(&lock3);  Lock(&lock1);  
Unlock(&lock3); Unlock(&lock1);
```



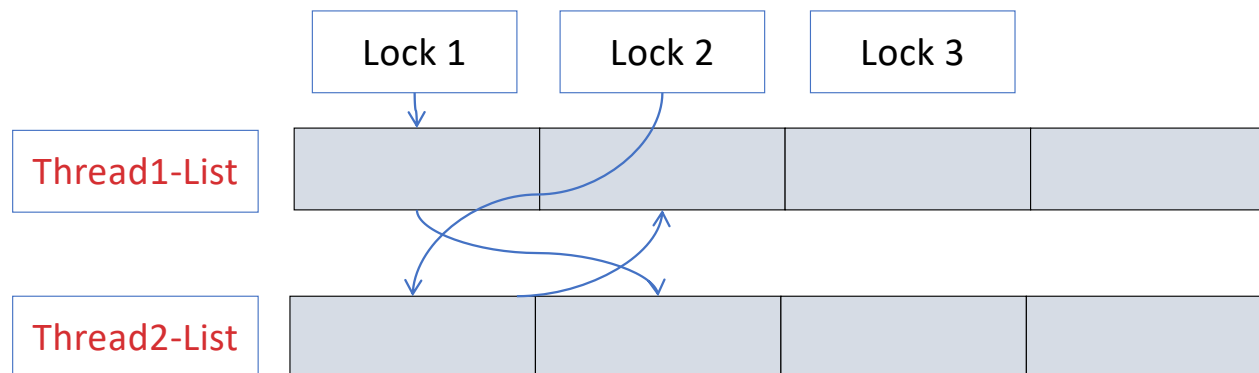
Recording Synchronizations

```
Thread1:      Thread2 :  
              Lock(&lock2);  
Lock(&lock1);  Unlock(&lock2);  
Lock(&lock2);  Lock(&lock1);  
Unlock(&lock2);  
Unlock(&lock1);  
Lock(&lock3);  Unlock(&lock1);  
Unlock(&lock3);
```



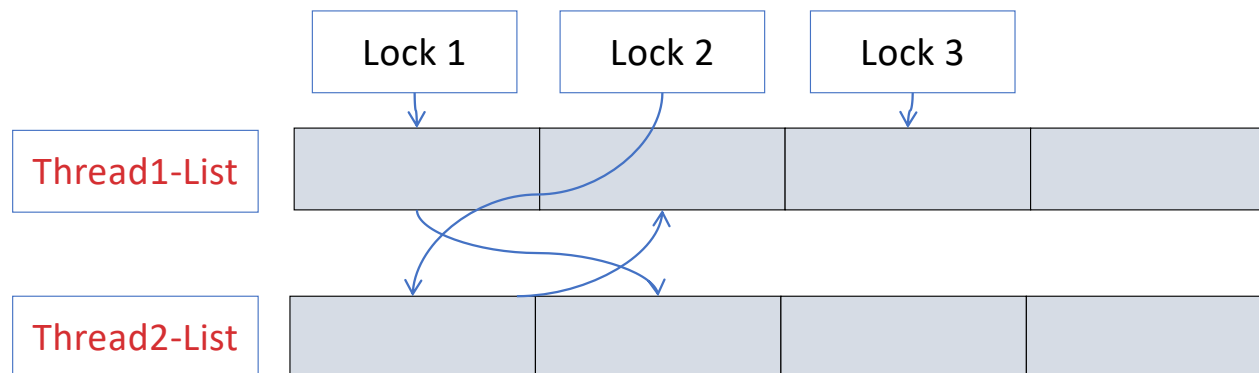
Recording Synchronizations

```
Thread1:      Thread2 :  
              Lock(&lock2);  
Lock(&lock1);  Unlock(&lock2);  
Lock(&lock2);  
Unlock(&lock2);  
Unlock(&lock1);  
Lock(&lock3);  
Unlock(&lock3); Lock(&lock1);  
                Unlock(&lock1);
```




Recording Synchronizations

```
Thread1:                                Thread2 :  
                                           Lock(&lock2);  
Lock(&lock1);                             Unlock(&lock2);  
Lock(&lock2);                               
Unlock(&lock2);                             
Unlock(&lock1);                             
Lock(&lock3);                             Lock(&lock1);  
Unlock(&lock3);                           Unlock(&lock1);
```



Benefits of Such Recording



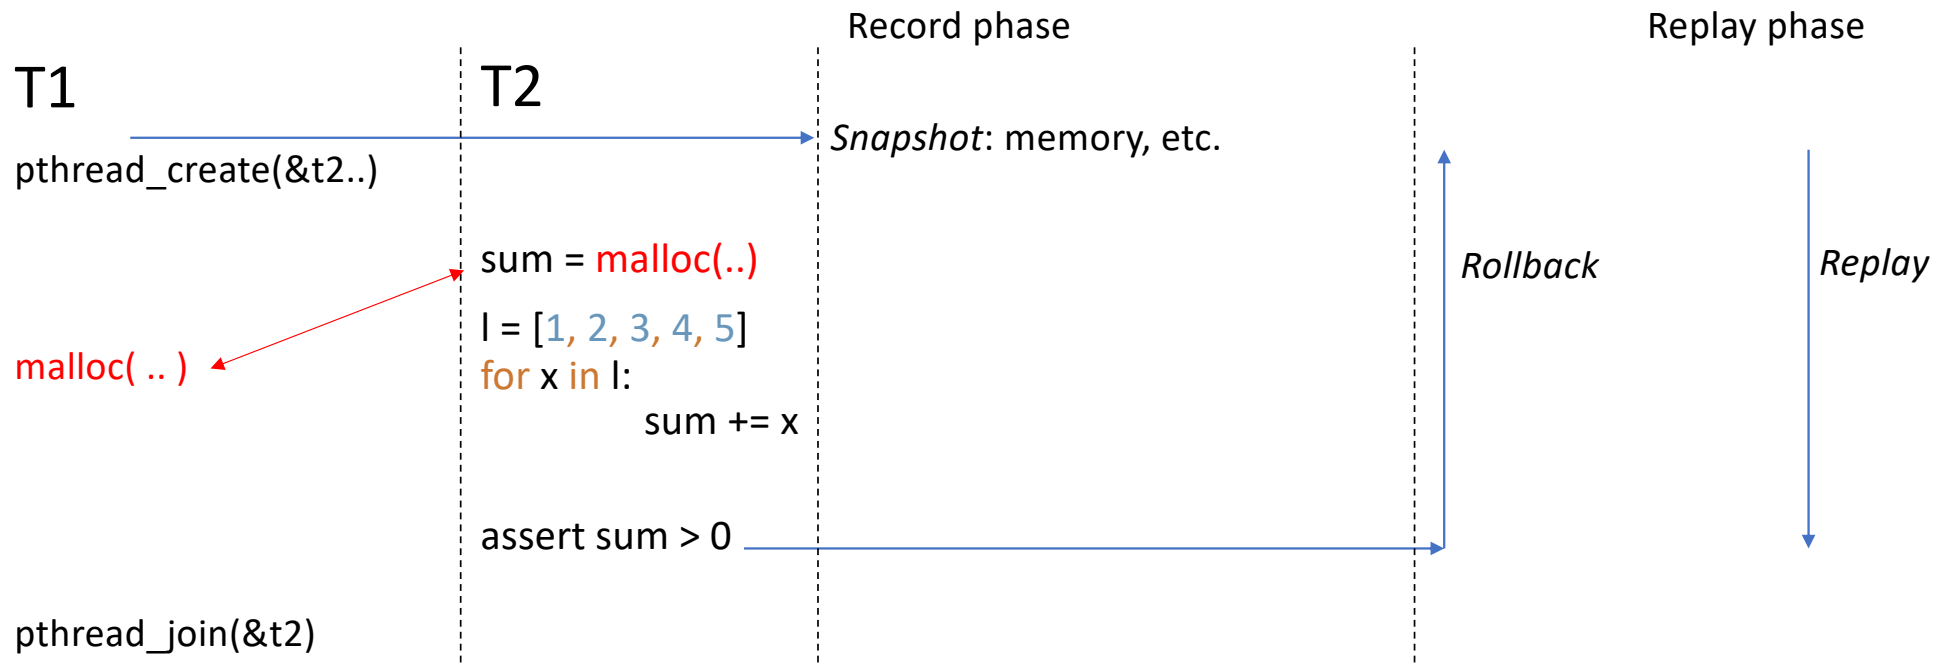
- Local-order recording guarantees identical reproduction
- Pre-allocated list avoids allocation overhead
- No additional locks required for recording
- Events are connected via per-thread or per-sync-variable lists, which can be used directly for reproduction

If an event is at the header of both per-thread list and its per-variable-list, the thread can proceed. Otherwise, wait.

Addressing sources of non-determinism

- Single thread: system calls
 - Repeatable, recordable, revocable, deferrable, irrevocable
- **Multithreading**
 - Thread lifecycle
 - Synchronization
 - **Memory allocators**
 - **Racy memory accesses**

Memory allocation



Identical Allocations/Deallocations


✧ Observation:

- Memory allocations/deallocations inside each thread is determined by the program order

✧ **Basic approach:**


- Every thread has its own heap
- Controlling each thread's interaction with other threads

Identical Allocations/Deallocations



- Allocation: deterministically fetch blocks via a global lock
Deallocation: freed objects are returned to the current thread's heap


Racy accesses



Observations

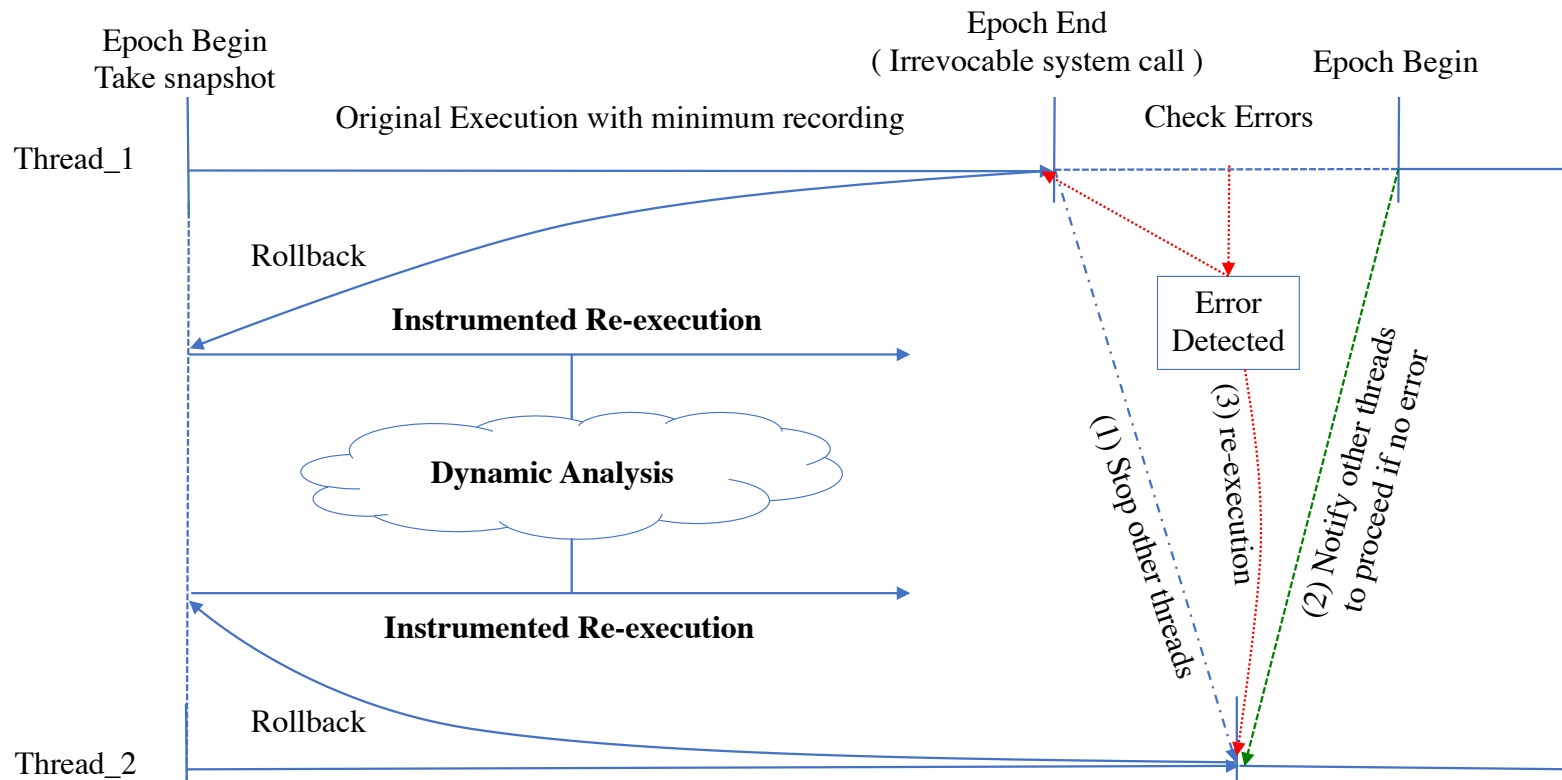
- Ordering every load/store will create significant overhead
- Unnecessary logging overhead if I need not rollback the epoch
- Most code is non-racy

Handling Races in Replays



If replay diverges, possibly caused by races,
iReplayer re-executes until an identical schedule is found!

Overview of iReplayer

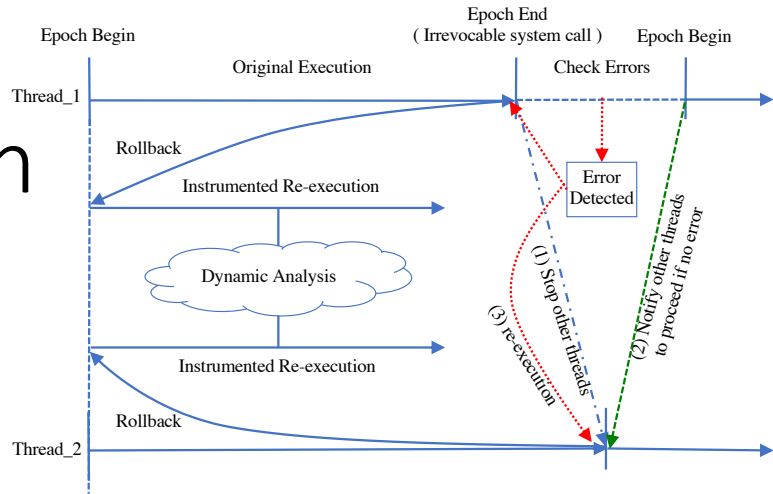


How to achieve in-situ and identical RnR efficiently?


Other Evaluation

- Identical Re-execution
 - All applications were identically reproduced
- Reproducing the race of Crasher
 - 99.8718%: in one execution
 - 0.1088%: in two executions
 - 0.0121%: in three executions
 - 0.0073%: \geq four executions

Conclusion



Category	Syscall Examples
Repeatable	getpid, getcwd
Recordable	gettimeofday, mmap, open
Revocable	file read/write
Deferrable	close, munmap, thread exits
Irrevocable	fork



Handling Synchronizations

